

# OpenMP - 1

CSE 625

Fall 2022

Computer Engineering and Computer Science  
University of Louisville

# Sources and References

[1] Official OpenMP Website

<http://www.openmp.org>

[2] Microsoft OpenMP Functions Reference

<https://docs.microsoft.com/en-us/cpp/parallel/openmp/reference/openmp-functions?view=vs-2019>


[3] VS 2019 OpenMPExamples Project

# Contents


- OpenMP Concepts
- Parallel Construct
- Work-Sharing Constructs

# What is OpenMP?

- An Application Program Interface (API) that may be used to explicitly direct *multi-threaded*, *shared memory* parallelism in Fortran and C/C++.  
(OpenMP is designed for multi-processor/core, shared memory machines. The underlying architecture can be shared memory UMA or NUMA. )



Uniform Memory Access



Non-Uniform Memory Access

# OpenMP History

- The OpenMP Architecture Review Board (ARB) published its first API specifications, OpenMP for Fortran 1.0, in October 1997. In October 1998, the OpenMP C/C++ specifications were released.
- OpenMP 2.0 for Fortran released in 2000  
OpenMP 2.0 for C/C++ released in 2002

## **OpenMP History - continued**

- OpenMP 2.5 Combined C/C++/Fortran released in 2005
- OpenMP 3.0 released in May, 2008
- OpenMP 4.0 released in July, 2014
- OpenMP 4.5 released in November 15, 2015
- OpenMP 5.0 released on November 8, 2018

# OpenMP API Summary

- Compiler Directives

For example:

```
#pragma omp parallel  
std::cout << "Hello, world.\n";
```

- Library Routines

For example

```
omp_get_thread_num()    // get thread id  
omp_set_num_threads (2) // using 2 threads
```

- Environment Variables

For example

```
setenv OMP_NUM_THREADS 8
```

# Program OpenMP using VS 2019

1. Create a new VS 2019 C++ console project
2. Open (Configuration) Property of the project (debug or release).
3. Click C/C++ → Language  
and modify the OpenMP Support as needed.
4. Include OpenMP header

```
#include <omp.h>
```

Visual C++ (2015-2019) supports only the OpenMP 2.0 standard.



# OpenMP Hello-World Example

## VS 2019 Project - OpenMPExamples


```
#include <omp.h>
#include <iostream>

void omp_hello_world()
{
    std::cout << "The CPU has " << omp_get_num_procs()
               << " cores. \n\n";

    #pragma omp parallel
    {
        int id = omp_get_thread_num();
        std::cout << "Hello, world greeting from thread "
                  << id << std::endl;
    }
}
```

# OpenMP Execution Model

Multi-Threading, fork-join model:




Using *#pragma omp parallel* directive

# OpenMP Execution Model - continued

- Defines parallel region over structured block of code. Threads are created as 'parallel' pragma.
- Data is shared among threads unless specified otherwise (note variables define in the block are private).

C/C++ :

```
#pragma omp parallel
{
    block
}
```



# OpenMP Execution Model - continued

- The master thread forks off a number of concurrent threads which execute structured blocks code in parallel.
- The number of concurrent threads:
  - default (number of cores, `omp_get_num_procs()`)
  - via a function call like this `omp_set_num_threads (4);`
  - via compiler directive like this  

```
#pragma omp parallel num_threads(4)
```
  - specified in an environment variable
- After the execution of the parallel structured block, the threads join into the master thread.
- Both data-parallel and task-parallel can be achieved.

# Thread ID

- Each thread running in parallel region has a unique id assigned to it. The unique thread id can be obtained by

```
int id = omp_get_thread_num();
```

- The unique ids are assigned as 0,1,2,3,4 ....
- Consider this example

```
#pragma omp parallel
{
    int id = omp_get_thread_num();
    std::cout << id << std::endl;
}
```

- The master thread has an id of 0.

# OpenMP Programming Ideas

- Create teams of threads for parallel execution.
- Specify how to share work among members of a thread team (load balancing).
- Declare *shared* and *private* variables as needed.
- Synchronize threads to enable them to perform certain operations exclusively without interferences from other threads (e.g. handling race conditions).

# Contents

- OpenMP Concepts
- **Parallel Construct**
- Work-Sharing Constructs

# Parallel Construct

## Syntax

```
#pragma omp parallel [clause ...] newline
    if (scalar_expression)
    private (list)
    shared (list)
    default (shared | none)
    firstprivate (list)
    reduction (operator: list)
    copyin (list)
    num_threads (integer-expression)
```

structured\_block



# Parallel Construct - continued

- One entry and one exit point
- There is an implied *barrier* at the end of a parallel region. Only the master thread continues execution past this point.
- How many threads? Use the default or the following clause:

`num_threads (integer-expression)`

For example:

```
#pragma omp parallel num_threads(4)
{
    int id = omp_get_thread_num();
    std::cout << id << std::endl;
} // Implicit barrier here
```

# Parallel Construct Example

## VS 2019 Project - OpenMPExamples

```
#include <omp.h>
#include <iostream>

void omp_hello_world_2()
{
    int nthreads, tid;

    #pragma omp parallel private(tid)
    {
        /* Obtain and print thread id */
        tid = omp_get_thread_num();
        printf("Hello World from thread = %d\n", tid);

        /* Only master thread does this */
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }
    } /* All threads join master thread and terminate */
}
```

# Contents

- OpenMP Concepts
- Parallel Construct
- Work-Sharing Constructs

# Sharing Work among Concurrent Threads

- Redundantly execute all of the structured block code (not very useful indeed).
- Execute on selected data items based on the thread id (like programming in POSIX pthreads or C++11 thread).
- Distribute the work in a Work-Sharing Construct, for example,

```
#pragma omp parallel
#pragma omp for
for (i = 0; i < _N; i++)
    sum += a[i];
```

How does OpenMP distribute for-loop work to running threads (job scheduling strategies)?

# Work-Sharing Constructs

- Loop construct

```
#pragma omp for [clause ...]
```

- Section work unit

```
#pragma omp sections [clause ...]
```

- Single

```
#pragma omp single [clause ...]
```

# Work-Sharing Constructs Concepts

- A work-sharing construct must be enclosed dynamically within a parallel region in order for the directive to execute in parallel.

```
void worker (int jobId)
{
    int tid = omp_get_thread_num();
    // #pragma omp critical
    printf("Thread %d does job %d\n", tid, jobId);
}
```

```
#pragma omp parallel num_threads(2)
{
    #pragma omp sections
    {
        #pragma omp section
        worker(10);

        #pragma omp section
        worker(20);
    }
}
```

# Work-Sharing Constructs Concepts -continued

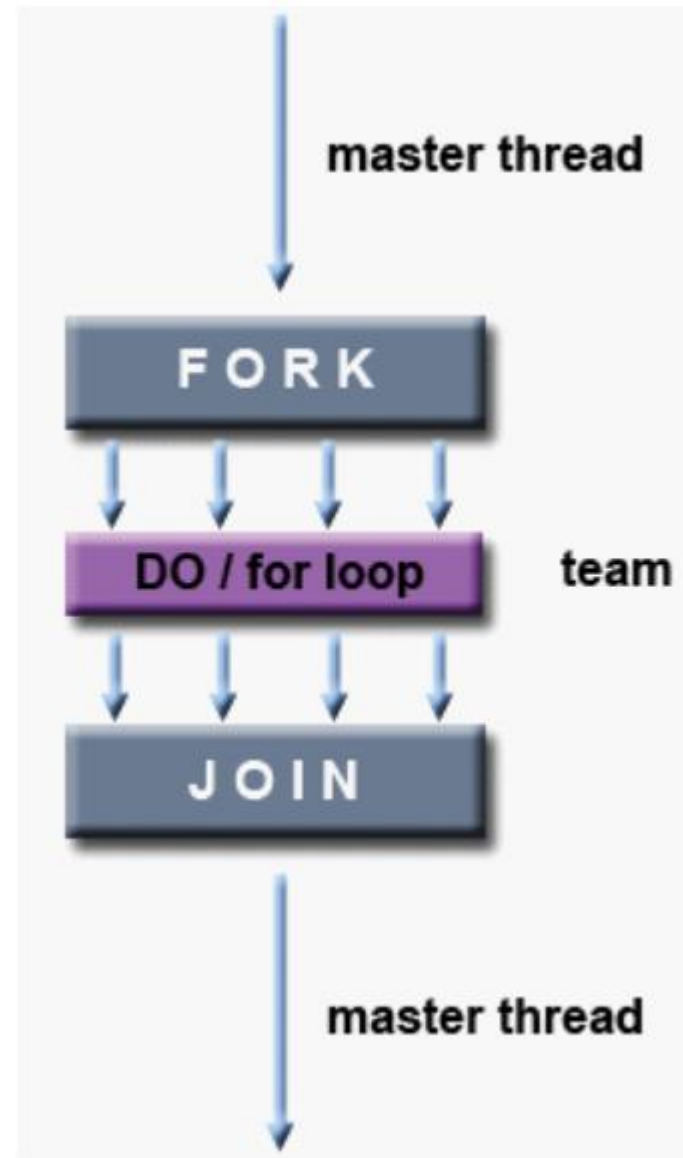
- A work-sharing construct does not launch new threads, has an implicit barrier at the end, but does not have an implicit barrier on entry.
- Threads wait at a barrier until the last thread has reached the barrier.

```
#pragma omp parallel
{
    #pragma omp for
    for(i = 0; i < N; i++)
        a[i] = b[i] + c[i];

    #pragma omp for
    for (i = 0; i < N; i++)
        sum += a[i];
}
```

# Work-Sharing Constructs Concepts - continued


- DO/for - shares iterations of a loop across the team. Represents a type of "data parallelism".






# Work-Sharing Constructs Concepts - continued

- **SECTIONS** - breaks work into separate, discrete sections. Each section is executed by a thread. Can be used to implement a type of "functional or task parallelism".



# Work-Sharing Constructs - continued

- **SINGLE** - serializes a section of code.



# Work-Sharing Constructs Concepts - continued

## Restrictions:

- A work-sharing construct must be enclosed dynamically within a parallel region in order for the directive to execute in parallel.
- Work-sharing constructs must be encountered by all members of a team or none at all.
- Successive work-sharing constructs must be encountered in the same order by all members of a team.

# For Work-Sharing Construct Syntax

```
#pragma omp for [clause ...] newline
                    schedule (type [,chunk])
                    ordered
                    private (list)
                    firstprivate (list)
                    lastprivate (list)
                    shared (list)
                    reduction (operator: list)
                    collapse (n)
                    nowait
```

for\_loop

# For Work Sharing Concepts


```
#pragma omp parallel
#pragma omp for
    for (I=0; I<N; I++) {
        Do_Work(I) ;
    }
```

- Splits loop iterations into threads.
- Must be in the parallel region,
- Must precede the loop

# For Work Sharing Concepts - continued

```
#pragma omp parallel num_threads(3)  
#pragma omp for  
    for(i = 1, i < 13, i++)  
        c[i] = a[i] + b[i];
```

- Threads are assigned an independent set of iterations
- Threads must wait at the end of work-sharing construct



# For Work Sharing Concepts - continued

These two code segments are equivalent:

```
#pragma omp parallel
{
    #pragma omp for
    for (i=0; i< MAX; i++) {
        res[i] = huge();
    }
}
```

```
#pragma omp parallel for
    for (i=0; i< MAX; i++) {
        res[i] = huge();
    }
```

# Matrix-Vector Multiplication

## VS 2019 Project - OpenMPExamples

```
#include <omp.h>
void omp_matrix_vector(
    std::vector<float>& A,
    std::vector<float>& x,
    std::vector<float>& b,
    int m,
    int n)
{
    #pragma omp parallel for
    // notice in OpenMP 2.0 array index must be int type
    for (int row = 0; row < m; row++)
    {
        float accum = float(0);
        for (int col = 0; col < n; col++)
            accum += A[row * n + col] * x[col];

        b[row] = accum;
    }
}
```



# For Work Sharing Example

```
#include <omp.h>
#define N 1000
#define CHUNKSIZE 100

main(int argc, char *argv[])
{
    int i, chunk;
    float a[N], b[N], c[N];

    /* Some initializations */
    for (i=0; i < N; i++)
        a[i] = b[i] = i * 1.0;
    chunk = CHUNKSIZE;

    #pragma omp parallel shared(a,b,c,chunk) private(i)
    {

        #pragma omp for schedule(dynamic,chunk) nowait
        for (i=0; i < N; i++)
            c[i] = a[i] + b[i];
    } /* end of parallel region */
}
```

# For Work Sharing Example - nowait

```
#pragma omp parallel default(none) shared(n,a,b) private(i)
{
    #pragma omp for nowait
    for (i = 0; i < n; i++)
    {
        printf("Thread %d executes first loop iteration %d\n",
               omp_get_thread_num(),i);
        a[i] = i;
    }

    #pragma omp for
    for (i = 0; i < n; i++)
    {
        printf("Thread %d executes second loop iteration %d\n",
               omp_get_thread_num(),i);
        b[i] = 2 * a[i];
        //b[i] = a[n - i - 1];
    }
}
```

Is the `nowait` clause correctly used?

# For Work Sharing **schedule** Clause

- **schedule**: describes how iterations of the loop are divided among the threads in the team. The default schedule is implementation dependent.
- **schedule clause syntax**

**schedule** (<kind> [, <chunk\_size>])

- **schedule** <kind> (i.e. algorithms to distribute work to threads):
  - static** (default)
  - dynamic**
  - guide**
  - runtime** (specified by the **OMP\_SCHEDULE** environment variable.)

## For Work Sharing Schedule - **static**

- For iterations are divided into equal size (specified by the `<chunk_size>`, or default to `#number_intertions/number_threads`).
- The chunks are assigned to the threads statically in a round-robin manners in the order the thread id.

```
#define N 1000
int i;
float a[N], b[N], c[N];
for (i=0; i < N; i++)
a[i] = b[i] = i * 1.0;

#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp for schedule(static, 10) nowait
    for (i = 0; i < N; i++)
        c[i] = a[i] + b[i];
}
```

## For Work Sharing Schedule - **dynamic**

- Each thread executes a chunk of iterations, then requests another chunk until the job (chunk) queue is empty.
- The size of chunks is specified by `chunk_size`, which is default to 1.

```
#define N 1000
int i;
float a[N], b[N], c[N];
for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;

#pragma omp parallel shared(a,b,c) private(i)
{
    #pragma omp for schedule(dynamic) nowait
    for (i = 0; i < N; i++)
        c[i] = a[i] + b[i];
}
```

# For Work Sharing Schedule - **guide**

- Similar to dynamic, but the `chunk_size` decreases in time, for example:

$$\text{chunk\_size} \propto (\text{number\_unsigned\_iterations} / \text{number\_threads})$$

- For `chunk_size`, of  $k$ ,  $k > 1$ , the size of chunks is not fewer than  $k$ .
- The `chunk_size` is default to 1.

# For Work Sharing **reduction** Clause

- Syntax

reduction (operator: list)

Operation	Fortran	C/C++	Initialization
Addition	+	+	0
Multiplication	*	*	1
Subtraction	-	-	0
Logical AND	.and.	&&	0
Logical OR	.or.		.false. / 0
AND bitwise	iand	&	all bits on / 1
OR bitwise	ior		0
Exclusive OR bitwise	ieor	^	0
Equivalent	.eqv.		.true.
Not Equivalent	.neqv.		.false.
Maximum	max	max	Most negative #
Minimum	min	min	Largest positive #

# For Work Sharing **reduction** Clause Example

```
#include <omp.h>
void main(int argc, char *argv[])
{
    int    i, n, chunk;
    float a[100], b[100], result;

    n = 100;
    chunk = 10;
    result = 0.0;
    for (i=0; i < n; i++) {
        a[i] = i * 1.0;
        b[i] = i * 2.0;
    }
    #pragma omp parallel for private(i) schedule(static,chunk) \
        reduction(+:result)
    for (i=0; i < n; i++)
        result = result + (a[i] * b[i]);
}
```



# Sections Work-Sharing Construct Syntax

```
#pragma omp sections [clause ...]  newline
    private (list)
    firstprivate (list)
    lastprivate (list)
    reduction (operator: list)
    nowait
{

    #pragma omp section  newline

    structured_block

    #pragma omp section  newline

    structured_block

}
```

# Sections Work-Sharing Construct Example

```
#include <omp.h>
#define N 1000

void main(int argc, char *argv[])
{

    int i;
    float a[N], b[N], c[N], d[N];

    /* Some initializations */
    for (i = 0; i < N; i++)
    {
        a[i] = i * 1.5;
        b[i] = i + 22.35;
    }
}
```

# Sections Work-Sharing Construct Example - continued

```
#pragma omp parallel shared(a,b,c,d) private(i)
{

    #pragma omp sections nowait
    {

        #pragma omp section
        for (i = 0; i < N; i++)
            c[i] = a[i] + b[i];

        #pragma omp section
        for (i = 0; i < N; i++)
            d[i] = a[i] * b[i];

    } /* end of sections */

} /* end of parallel region */

} /* end of main */
```

# Single Work-Sharing Construct Syntax

```
#pragma omp single [clause ...] newline  
    private (list)  
    firstprivate (list)  
    nowait
```

structured\_block

# Single Work-Sharing Example

```
#include <omp.h>
#include <iostream>

#define _N (10)
void singleDemo()
{
    int a, b[_N], i;

    #pragma omp parallel shared(a, b) private(i)
    {
        #pragma omp single
        {
            a = 10;
            std::cout << "Single construct "
                      " executed by thread "
                      << omp_get_thread_num()
                      << std::endl;
        }
    }
}
```

# Single Work-Sharing Example - continued

```
// end of #pragma omp single

// implicit barrier here

#pragma omp for
for (i = 0; i < _N; i++)
    b[i] = a;

} // end of #pragma omp parallel

for (i = 0; i < _N; i++)
    printf ("b[%d] = %d\n", i, b[i]);

} // end of singleDemo function
```

# One More Work-Sharing Example

```
#include <omp.h>
#define N      1000
#define CHUNKSIZE  100

void main(int argc, char *argv[])
{
    int i, chunk;
    float a[N], b[N], c[N];

    /* Some initializations */
    for (i=0; i < N; i++)
        a[i] = b[i] = i * 1.0;
    chunk = CHUNKSIZE;

    #pragma omp parallel for \
        shared(a,b,c,chunk) private(i) \
        schedule(static,chunk)
    for (i=0; i < N; i++)
        c[i] = a[i] + b[i];
}
```