

CSE 625 Parallel Programming
Project 3
By Caleb Klenda

Machine Specifications

The project was all performed on my home computer with the following specifications:

- **CPU**
Intel(R) Core(TM) i9-10900K CPU @ 3.70GHz
AVX2 (256-bit MM registers)
10 cores / 20 threads
20 MB Intel Smart Cache (L3-cache)
- **RAM**
32 GB DDR4 RAM
- **GPU**
TUF RTX3080 (Ampere GPU)
8704 CUDA cores
5 MB of L2-Cache
10GB GDDR6X

Problem 1

```
[1] import numpy as np
import matplotlib.pyplot as plt
✓ 4.4s Python
```

```
[2] mis_match = [(115, 8111), (195, 47020), (241, 9732), (268, 47938), (300, 49308), (320, 33406), (321, 10485), (341, 34266), (358, 19138), (381,
✓ 0.6s Python
```

```
[3] len(mis_match)
✓ 0.4s Python
```

... 309

Read in Data

+ Code + Markdown

```
[4] # Read train images (60,000x28x28 float32)
train_images = np.fromfile('./data/train-images.bin', dtype = np.float32)
train_images = train_images.reshape(60000, 28, 28)

# Read train images (60,000 ubyte)
train_labels = np.fromfile('./data/train-labels.bin', dtype = np.ubyte)

# Read test images (10,000x28x28 float32)
test_images = np.fromfile('./data/test-images.bin', dtype = np.float32)
test_images = test_images.reshape(10000, 28, 28)

# Read test images (10,000 ubyte)
test_labels = np.fromfile('./data/test-labels.bin', dtype = np.ubyte)
✓ 0.5s Python
```

Plot

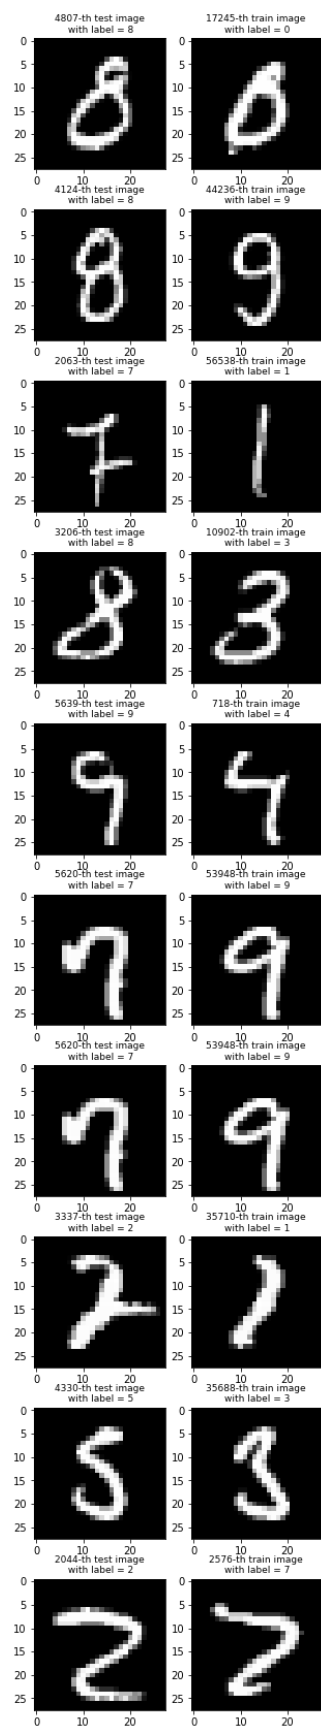
```
#utility for plotting
def pairPlot(indexes, i):
    test_idx, train_idx = indexes
    plt.subplot(10, 2, (2*i)+1)
    plt.imshow(test_images[test_idx], cmap='gray')
    plt.title(str(test_idx) + "-th test image\n with label = " + str(test_labels[test_idx]), fontsize = 9)
    plt.subplot(10, 2, (2*i)+2)
    plt.imshow(train_images[train_idx], cmap='gray')
    plt.title(str(train_idx) + "-th train image \nwith label = " + str(train_labels[train_idx]), fontsize = 9)

def pickandPlotRandomPairs(numberOfPairs):
    plt.figure(figsize=(5,30)) # width, height in inch of the plot
    for i in range(0, numberOfPairs):
        randomPair = mis_match[np.random.choice([k for k in range(0,len(mis_match))])]
        pairPlot(randomPair, i)

pickandPlotRandomPairs(10)
```

[*] ✓ 1.9s Python

... 48074th test image with label = 8 17245-th train image with label = 0



Mis_match.ipynb
CSE625H3-20. 200
Dataset.ipynb

mismatch_klenda.ipynb > M*Plot > #usi Enter an index between 0-308: (Press 'Enter' to confirm or 'Escape' to cancel) choice = int(
+ Code + Markdown ▶ Run All ⌵ Clear Outputs of All Cells ↺ Go To Running Cell ⏹ Restart ⏏ Interrupt 📄 Variables 📄 Outline ...



#using utility from before
def pairAndPlotUser():
 choice = -1
 while(choice < 0 or choice > 308):
 choice = int(input('Enter an index between 0-308: '))
 if choice > 0 and choice < 308:
 print("You chose: ", mis_match[choice])
 pairPlot(mis_match[choice], 0)

plt.figure(figsize=(4,14))
pairAndPlotUser()

[9] 25.2s

#using utility from before
def pairAndPlotUser():
 choice = -1
 while(choice < 0 or choice > 308):
 choice = int(input('Enter an index between 0-308: '))
 if choice > 0 and choice < 308:
 print("You chose: ", mis_match[choice])
 pairPlot(mis_match[choice], 0)

plt.figure(figsize=(4,14))
pairAndPlotUser()

[9] ✓ 31.9s

... You chose: (4435, 19434)

4435-th test image with label = 3
19434-th train image with label = 7



Problem 2

In the CodeBlocks project, `All_Pair_distance`, it implements four functions to compute the pair-wise distance matrix of MNIST train images (loaded from `train-images.bin`). These four methods are:

- 1- `sequential_all_pairs` (sequential computing)
- 2- `block_all_pairs` (C++ multi threads - block work distribution)
- 3- `block_cyclic_all_pairs` (C++ multi threads - block cyclic work distribution)
- 4- `dynamic_all_pairs` (C++ multi threads - dynamic work distribution)

2.1

Matrix Size	400	800	10,000	20,000	30,000	60,000
Method 1	0.0503091	0.200686	32.9804	136.509	322.867	1560.27
Method 2 12 threads	0.0086379	0.0312908	6.41424	34.1735	97.6546	425.77
Method 3 12 threads Chunk size 2	0.0052788	0.0184834	2.78006	11.524	28.0694	112.191
Method 4 12 threads Chunk size 2	0.0050414	0.0183471	2.77072	11.5129	27.8863	111.889

2.2 (8 points) In the report, explain the key ideas of the function, `dynamic_all_pairs`, of its work distribution implementation and how and why the `std::mutex` object is used.

`dynamic_all_pairs` assign chunks to threads at runtime allowing it to adapt the the problem it is solving. `global_lower` which allows access to the first row of the currently processed chunk, which, whenever a thread runs out of work, it will reference to determine what it should do next. Because multiple threads are accessing and modifying `global_lower`, a mutex is needed so that only one thread can use that resource at a time

Problem 3

Work amount (i.e., the number of outmost iterations) done by each thread for 2 threads on `block_all_pairs` for $m=60,000$ mxm matrix

We know that for `block_all_pairs` $T(i) = i + 1$. When using three threads, we split the work into the following:

$$\begin{aligned} W(1) &= \sum_{i=0}^{\frac{m}{3}-1} T(i) = \sum_{i=0}^{\frac{m}{3}-1} (i + 1) = \sum_{i=0}^{\frac{m}{3}-1} (i) + \frac{m}{3} = \frac{\left(\frac{m}{3} - 1\right) \left(\frac{m}{3}\right)}{2} + \frac{m}{3} = \frac{m^2}{18} - \frac{m}{6} + \frac{m}{3} \\ &= \frac{m^2}{18} + \frac{m}{6} \end{aligned}$$

$$\begin{aligned} W(2) &= \sum_{i=\frac{m}{3}}^{\frac{2m}{3}-1} T(i) = \sum_{i=\frac{m}{3}}^{\frac{2m}{3}-1} (i + 1) = \sum_{i=\frac{m}{3}}^{\frac{2m}{3}-1} (i) + \frac{m}{3} = \sum_{i=0}^{\frac{m}{3}-1} \left(\frac{m}{3} + i\right) + \frac{m}{3} = \\ &= \sum_{i=0}^{\frac{m}{3}-1} \left(\frac{m}{3}\right) + \sum_{i=0}^{\frac{m}{3}-1} (i) + \frac{m}{3} = \frac{m}{3} \left(\frac{m}{3}\right) + \frac{\left(\frac{m}{3} - 1\right) \left(\frac{m}{3}\right)}{2} + \frac{m}{3} \\ &= \frac{m^2}{9} + \frac{\frac{m^2}{9} - \left(\frac{m}{3}\right)}{2} + \frac{m}{3} = \frac{m^2}{9} + \frac{m^2}{18} - \frac{m}{6} + \frac{m}{3} = \frac{m^2}{6} + \frac{m}{6} \end{aligned}$$

$$W(3) = \sum_{i=\frac{2m}{3}}^{m-1} T(i) = \sum_{i=\frac{2m}{3}}^{m-1} (i + 1) = \sum_{i=\frac{2m}{3}}^{m-1} (i) + \frac{m}{3} = \sum_{i=0}^{\frac{m}{3}-1} \left(\frac{2m}{3} + i\right) + \frac{m}{3} =$$

$$\begin{aligned}
&= \sum_{i=0}^{\frac{m}{3}-1} \binom{2m}{3} + \sum_{i=0}^{\frac{m}{3}-1} (i) + \frac{m}{3} = \frac{m}{3} \binom{2m}{3} + \frac{\left(\frac{m}{3}-1\right)\left(\frac{m}{3}\right)}{2} + \frac{m}{3} \\
&= \frac{2m^2}{9} + \frac{\frac{m^2}{9} - \left(\frac{m}{3}\right)}{2} + \frac{m}{3} = \frac{2m^2}{9} + \frac{m^2}{18} - \frac{m}{6} + \frac{m}{3} = \frac{m^2}{2} + \frac{m}{6} \\
&= \frac{5m^2}{18} + \frac{m}{6}
\end{aligned}$$

So, in the general case the threads work as follows:

$$W(1) = \frac{m^2}{18} + \frac{m}{6}, W(2) = \frac{m^2}{6} + \frac{m}{6}, W(3) = \frac{5m^2}{18} + \frac{m}{6}$$

- We plug in 60,000 for m to achieve $W(1)_{60000} = 2,000,010,000$
- We plug in 60,000 for m to achieve $W(2)_{60000} = 600,010,000$
- We plug in 60,000 for m to achieve $W(3)_{60000} = 1,000,010,000$

From this we can observe that as $m \rightarrow \infty$ that $W(1)$ accounts for 11% of the work, $W(2)$ accounts for 33% of the work, and the $W(3)$ accounts for the remaining 55% of the work.

Problem 4

Work amount (i.e., the number of outmost iterations) done by each thread for 2 threads on `block_cyclic_all_pairs`,

Using the formula given $F(i0, c) = a \left(\frac{c^2 + 2ic + c}{2} \right)$ we compute the following for each thread:

- $W(1) = 1 \left(\frac{\left(\frac{m}{3}\right)^2 + \frac{2(0)m}{3} + \frac{m}{3}}{2} \right) = \frac{m^2}{18} + \frac{m}{6} = \frac{\frac{1}{3}m^2 + m}{6}$
- $W(2) = 1 \left(\frac{\left(\frac{m}{3}\right)^2 + \frac{2\left(\frac{m}{3}\right)m}{3} + \frac{m}{3}}{2} \right) = \frac{m^2}{18} + \frac{2m^2}{18} + \frac{m}{6} = \frac{m^2 + m}{6}$

- $$W(3) = 1 \left(\frac{\left(\frac{m}{3}\right)^2 + \frac{2\left(\frac{2m}{3}\right)m}{3} + \frac{m}{3}}{2} \right) = \frac{m^2}{18} + \frac{4m^2}{18} + \frac{m}{6} = \frac{5m^2}{18} + \frac{m}{6}$$

We can see that these exactly match the results from Problem 3 so:

- We plug in 60,000 for m to achieve $W(1)_{60000} = 2,000,010,000$
- We plug in 60,000 for m to achieve $W(2)_{60000} = 600,010,000$
- We plug in 60,000 for m to achieve $W(3)_{60000} = 1,000,010,000$